

System Security (201700086)

Secure Systems Engineering #2 - Add seccomp rules

Æde Symen Hoekstra, Luigi Coniglio (Group 24)

June 18, 2018

1 Syscall employed

In order to obtain the list of syscalls employed by the child process responsible of handling the query we traced the execution using the command *strace*:

```
$ strace -f ./randomcatserver
```

The listing below shows an example of output of the above command:

```
1 [pid 16266] close(3) = 0
2 [pid 16218] close(4) = 0
3 [pid 16218] accept(3, <unfinished ...>
4 [pid 16266] read(4, "G", 1) = 1
5 [pid 16266] read(4, "E", 1) = 1
6 [pid 16266] read(4, "T", 1) = 1
7 ...
8 [pid 16266] read(4, "\n", 1) = 1
9 [pid 16266] write(1, "Random cat\n", 11) = 11
10 [pid 16266] write(4, "HTTP/1.0 200 OK\r\nContent-Type: i"... , 64) = 64
11 [pid 16266] open("images", O_RDONLY) = 3
12 [pid 16266] fstat(3, {st_mode=S_IFDIR|0755, st_size=4096, ...}) = 0
13 [pid 16266] fcntl(3, F_GETFL) = 0x8000 (flags O_RDONLY|O_LARGEFILE)
14 [pid 16266] fcntl(3, F_SETFD, FD_CLOEXEC) = 0
15 [pid 16266] brk(0) = 0x169f000
16 [pid 16266] brk(0x16c8000) = 0x16c8000
17 [pid 16266] getdents(3, /* 23 entries */, 32768) = 1104
18 [pid 16266] getdents(3, /* 0 entries */, 32768) = 0
19 [pid 16266] openat(3, "pexels-photo-326875.jpeg", O_RDONLY) = 5
20 [pid 16266] brk(0x16c0000) = 0x16c0000
21 [pid 16266] close(3) = 0
22 [pid 16266] read(5, "\377\330\377\340\0\20JFIF\0\1\1"... , 4096) = 4096
23 [pid 16266] write(4, "\377\330\377\340\0\20JFIF\0\1\1"... , 4096) = 4096
24 ...
25 [pid 16266] read(5, "\224\17\367/\&\367\30\366\253vWV"... , 4096) = 1662
```

```

26 [pid 16266] write(4, "\224\17\367/\&\367\30\366\253vWV"... , 1662) = 1662
27 [pid 16266] read(5, "", 4096) = 0
28 [pid 16266] shutdown(4, SHUT_WR) = 0
29 [pid 16266] close(4) = 0
30 [pid 16266] exit_group(0) = ?
31 [pid 16266] +++ exited with 0 +++

```

Those are the syscalls executed by the child process serving a GET request for the resource */randomcat* .

As we can see there server uses the *read* and *write* syscalls to communicate with the client and read the image, the syscall *open* to open the images directory and the cat image, the syscalls *fstat* and *fcntl* to retrieve some informations about the images directory, the syscall *brk* to allocate some memory, the syscall *getdents* to get what files are inside the images directory, *openat* to open the image file, *close* to close the file descriptors not longer needed, *shutdown* to shut down the socket and *exit_group* to terminate the execution.

2 Adding seccomp rules

In order to stop the process from using different syscall than those specified in the previous section, we execute the following function before the core operations:

```

1 int setup_seccomp() {
2     scmp_filter_ctx ctx;
3     ctx = seccomp_init(SCMP_ACT_KILL);
4     int calls[] = {SCMP_SYS(close),      SCMP_SYS(read),      SCMP_SYS(write),
5                   SCMP_SYS(open),      SCMP_SYS(fstat),      SCMP_SYS(fcntl),
6                   SCMP_SYS(brk),       SCMP_SYS(getdents),    SCMP_SYS(openat),
7                   SCMP_SYS(shutdown), SCMP_SYS(exit_group)};
8     int calls_length = sizeof(calls) / sizeof(calls[0]);
9     int i;
10
11     if (ctx == NULL) {
12         return -1;
13     }
14
15     for (i = 0; i < calls_length; i++) {
16         if (seccomp_rule_add(ctx, SCMP_ACT_ALLOW, calls[i], 0) < 0) {
17             seccomp_release(ctx);
18             fprintf(stderr, "adding rule %d failed \n", i);
19             return -1;
20         }
21     }
22
23     if (seccomp_load(ctx) < 0) {
24         seccomp_release(ctx);
25         fprintf(stderr, "loading seccomp failed\n");

```

```
26         return -1;
27     }
28     return 0;
29 }
```
